

Transformer-Based Belief Modeling for Offline Policy Learning in Partially Observed, Simultaneous-Action Games

Ethan [Author Name]

July 6, 2026

Abstract

Many real decision problems are partially observed, require simultaneous action commitment, and exhibit combinatorial structure that causes naive learners to memorize frequent patterns rather than infer hidden state. This thesis studies whether causal transformer sequence models can maintain useful beliefs over hidden opponent attributes from interaction history and whether explicit belief supervision improves offline policy learning under combinatorial distribution shift.

We propose a hierarchical tokenization scheme with within-turn self-attention and cross-turn causal attention, auxiliary belief prediction heads over opponent party slots, and a combinatorial evaluation protocol based on held-out team keys and move-set combinations. Experiments use large-scale human competitive game trajectories reconstructed from spectator replays. The scientific contribution is about belief inference and generalization, not about game-specific heuristics or lookahead search at inference time.

Keywords: partial observability, offline reinforcement learning, transformers, opponent modeling, combinatorial generalization, behavioral cloning

Acknowledgments

I thank my thesis advisor and the research community behind Metamon, poke-env, and the Pokemon Showdown ecosystem for open tools and datasets that make this work possible.

Contents

Acknowledgments	1
1 Introduction	7
1.1 Motivation	7
1.2 Problem Statement	7
1.3 Research Questions	7
1.4 Contributions	8
1.5 Evaluation Domain	8
1.6 Thesis Organization	8
2 Background	9
2.1 Partially Observable Decision Making	9
2.2 Simultaneous-Action Games	9
2.3 Transformers for Sequential Decision Making	9
2.4 Offline Reinforcement Learning	9
2.5 Combinatorial Generalization	10
2.6 Positioning Relative to Metamon	10
3 Benchmark Environment	11
3.1 Why Competitive Pokémon	11
3.2 Infrastructure	11
3.3 Scope of This Chapter	11
3.4 Action Space	11
3.5 Data Pipeline Status	11
4 Representation and Belief Targets	13
4.1 Design Goals	13
4.2 Thirteen-Token Turn Layout	13
4.3 History Window	13
4.4 Belief Supervision Targets	14
4.5 Combinatorial Keys for Evaluation	14
4.6 Implementation	14
5 Model and Training	15
5.1 Overview	15
5.2 Hierarchical Encoder	15
5.2.1 Within-Turn Self-Attention	15
5.2.2 Cross-Turn Causal Attention	15
5.3 Policy Head	15
5.4 Belief Heads	16
5.5 Belief Fusion Variant	16
5.6 Training Procedure	16

5.6.1	Stage A: Belief-Augmented Behavioral Cloning	16
5.6.2	Stage B: Conservative Offline RL Fine-Tuning	16
5.6.3	Stage C: Evaluation	16
5.7	Baselines	16
5.8	Architecture Summary	16
6	Experimental Setup	18
6.1	Dataset	18
6.2	Data Splits	18
6.2.1	Split A: Random Battle Split	18
6.2.2	Split B: Held-Out Team Cores	18
6.2.3	Split C: Held-Out Set Combinations	18
6.2.4	Split D: Temporal Split (Optional)	18
6.3	Belief Metrics	18
6.4	Policy Metrics	19
6.5	Ablations	19
6.6	Statistical Reporting	19
6.7	Minimum Viable Result	19
6.8	Compute and Reproducibility	19
7	Results	20
7.1	Overview	20
7.2	Phase 1 Baseline (Full Subset)	20
7.3	Policy Quality (RQ2)	20
7.3.1	Random Battle Holdout (Split A)	20
7.3.2	Live Evaluation	21
7.4	Belief Inference (RQ1)	21
7.4.1	Active-Opponent Move Belief	21
7.4.2	Species and Set-Combination Beliefs	21
7.5	Combinatorial Generalization (RQ3)	21
7.5.1	Held-Out Opponent Teams (Split B)	21
7.6	Architecture Ablations (RQ4)	22
7.6.1	Belief Fusion	22
7.7	Stage B Offline RL	22
7.8	Summary	23
7.8.1	Relation to Hypotheses	23
8	Analysis and Interpretability	24
8.1	Overview	24
8.2	Scaling Belief Versus Scaling Policy	24
8.3	Why the Team Generalization Gap Collapsed	24
8.4	Linking Beliefs to Actions	25
8.5	Implicit Belief in the Turn-Only Trunk	25
8.6	Gen 9 OU and Belief Target Selection	26
8.7	Decision Tree and Attention Studies	26
8.8	Takeaways for RQ1–RQ4	26
9	Limitations and Future Work	27
9.1	Domain Specificity	27
9.2	Label and Reconstruction Noise	27
9.3	Belief Labels After Reveal	27
9.4	Compute and Model Scale	27

9.5	Offline RL Stability	27
9.6	Gen 9 Team Preview	27
9.7	Future Directions	28
10	Conclusion	29
10.1	Contributions	29
10.2	Summary of Findings	29
10.3	Hypotheses Revisited	30
10.4	Broader Impact	30
10.5	Future Work	30
A	Implementation Reference	32
A.1	Repository Layout	32
A.2	Building the PDF	32
A.3	Training Commands	32
A.4	TurnTransformer Forward Pass	33

List of Figures

List of Tables

4.1	Turn token layout. Indices are zero-based.	13
5.1	Default hyperparameters for the belief transformer.	17
6.1	Planned ablation studies.	19
7.1	Top-1 action match on random validation battles (Split A), belief corpus, 2 000 battles, 20 epochs.	20
7.2	Move-slot belief accuracy on the validation split (2 000-battle corpus).	21
7.3	Action match under held-out opponent team keys (Split B).	22
7.4	Results summary on the 2 000-battle belief corpus (20 epochs).	23
8.1	Move-slot belief accuracy on validation battles (same split).	25
A.1	Key paths in the thesis codebase.	32

Chapter 1

Introduction

1.1 Motivation

Many decision problems in security, negotiation, finance, and multi-agent interaction share three properties that standard fully observed reinforcement learning benchmarks underrepresent.

Partial observability. Important state variables remain hidden until late in an interaction. An agent must infer latent structure from noisy clues rather than read full state directly.

Simultaneous action selection. Both sides commit to actions before the environment resolves the step. This differs from chess-like turn-taking and creates a distinct learning problem around anticipating opponent intent under uncertainty.

Combinatorial structure. The space of valid configurations grows quickly. A model that memorizes frequent patterns may appear strong on average yet fail on rare but valid combinations held out from training.

Transformers are a natural fit for these settings because they aggregate evidence across long interaction histories with causal attention. Offline learning is appropriate when high-quality human trajectories are available at scale and the deployment goal is a single forward pass without lookahead search or simulator queries at inference time.

1.2 Problem Statement

Given a sequence of first-person observations from past turns, we seek a model that

- (i) maintains a useful belief over hidden opponent attributes,
- (ii) uses that belief to choose strong actions in new interactions, and
- (iii) generalizes to opponent configurations outside the training distribution.

Formally, each turn produces an observation o_t derived from public state. Hidden variables z_t include unrevealed species, moves, items, abilities, and other private attributes. Actions a_t are chosen under simultaneous-move constraints. The model receives offline trajectories $\{(o_{1:T}, a_{1:T}, z_{1:T})\}$ reconstructed from full logs where z is known post hoc but not visible to the agent at decision time.

The central thesis question is whether explicit belief supervision and transformer-based history modeling improve both belief accuracy and policy performance under combinatorial distribution shift.

1.3 Research Questions

RQ1 (Belief inference) Can a causal transformer predict hidden opponent attributes more accurately than classical and neural baselines?

- RQ2 (Belief usefulness)** Does improved belief accuracy translate into better action selection under behavioral cloning and offline policy metrics?
- RQ3 (Combinatorial generalization)** Under held-out team and set splits, do belief-augmented policies degrade more gracefully than action-only models?
- RQ4 (Architecture)** How much gain comes from within-turn structure, cross-turn causal attention, and auxiliary belief heads versus flat sequence baselines?

1.4 Contributions

This thesis makes four contributions.

1. A structured tokenization and belief supervision framework for partially observed simultaneous-action games.
2. Empirical evidence on whether transformer history models learn calibratable beliefs over hidden opponent attributes from offline human play.
3. A combinatorial generalization benchmark protocol with team-key and set-combination held-out splits tied to belief and policy metrics.
4. An open-source implementation built on Metamon, poke-env, and Gymnasium compatible tooling.

1.5 Evaluation Domain

We evaluate on competitive Pokémon battles reconstructed from human spectator replays via the Metamon dataset [Grigsby et al., 2023, 2025]. Pokémon is used only as a benchmark. It combines partial observability, simultaneous move commitment, and combinatorial team structure in one domain with large-scale offline data. The methods apply to any domain with similar structure.

1.6 Thesis Organization

Chapter 2 reviews POMDPs, transformers in decision making, and offline RL. Chapter 3 briefly describes the benchmark environment. Chapters 4 and 5 present tokenization and the belief transformer. Chapters 6 to 8 cover the experiment plan, results, and interpretability analysis. Chapters 9 and 10 discuss limitations and future work.

Chapter 2

Background

2.1 Partially Observable Decision Making

A partially observable Markov decision process (POMDP) extends the MDP formalism with a latent state s_t that is not fully observed [Kaelbling et al., 1998]. The agent receives observations o_t and must maintain a belief $b_t(s) = P(s_t = s \mid o_{1:t}, a_{1:t-1})$ over latent state. Optimal control in POMDPs requires planning over belief space, which is often intractable in large discrete domains.

Recent work amortizes belief inference with neural sequence models that map interaction history to action distributions or latent embeddings [Chen et al., 2021, Liu et al., 2024]. This thesis asks whether structured auxiliary supervision over interpretable hidden attributes improves both belief quality and downstream policy robustness.

2.2 Simultaneous-Action Games

In simultaneous-move settings both agents select actions before the environment advances. The effective decision problem is a partially observed stochastic game. Unlike turn-based games, the agent cannot observe the opponent’s current action before committing. Policies must therefore encode predictions about opponent intent from history alone.

2.3 Transformers for Sequential Decision Making

The transformer architecture uses self-attention to relate tokens within a sequence [Vaswani et al., 2017]. Decision Transformer and related methods apply causal transformers to offline RL by modeling return-conditioned action sequences [Chen et al., 2021]. In competitive settings, opponent modeling methods such as TAO embed opponent trajectories for in-context adaptation [Liu et al., 2024]. The Self-Confirming Transformer generates fictitious belief tokens representing inferred opponent actions and feeds them back into policy decoding [Zhang et al., 2024].

Our approach differs in three ways. We predict structured hidden attributes per opponent slot rather than only next-opponent-action embeddings. We use a hierarchical within-turn and cross-turn encoder rather than a flat token stream. We evaluate combinatorial held-out team and set splits as a primary metric rather than leaderboard rank alone.

2.4 Offline Reinforcement Learning

Offline RL learns from a fixed dataset without online environment interaction [Sutton and Barto, 2018]. Behavioral cloning (BC) imitates expert actions and provides a stable initialization but suffers from covariate shift. Conservative Q-Learning (CQL) penalizes overestimation on

out-of-distribution actions [Kumar et al., 2020]. AMAGO provides another sequence-model fine-tuning path compatible with transformer backbones.

We use BC with auxiliary belief losses as Stage A and optional CQL or AMAGO fine-tuning as Stage B. Inference remains a single forward pass with no search.

2.5 Combinatorial Generalization

Humans generalize to novel combinations of known primitives [Lake et al., 2018]. Machine learning systems often fail when test combinations were absent from training even when individual primitives were seen. We adopt held-out team keys and held-out move-set tuples as explicit tests of combinatorial generalization rather than random battle splits alone.

2.6 Positioning Relative to Metamon

Grigsby et al. [2025] demonstrated that large sequence models trained on human Pokémon replays can reach strong play without search. That work establishes feasibility of transformer policies in this domain. Our thesis does not aim to beat their agents on ladder rank. Instead we study explicit belief inference, calibration, and OOD degradation under combinatorial splits, which their paper does not center.

Chapter 3

Benchmark Environment

3.1 Why Competitive Pokémon

Competitive Pokémon Singles (CPS) is a turn-based strategy game with simultaneous move selection each turn. Players choose teams of six Pokémon with moves, items, abilities, and in Generation 9 a terastallization type. At battle start much of the opponent roster is hidden. Species, sets, and intent must be inferred from damage ranges, switch patterns, item triggers, and move reveals.

CPS therefore combines partial observability, simultaneous commitment, long horizons (often 20–100+ turns), and combinatorial team structure. Human replay data spans more than a decade and grows daily on public ladders.

3.2 Infrastructure

We use three open components.

- **poke-env** bridges Python agents to a local Pokemon Showdown server for live evaluation.
- **Metamon** reconstructs first-person trajectories from spectator replays with full-state annotations for belief labels [Grigsby et al., 2023].
- **Gymnasium** provides the environment interface standard for RL tooling in this project.

3.3 Scope of This Chapter

This thesis is not about Pokémon mechanics, tier lists, or heuristic engines. We use Gen 9 OU human battles as the primary corpus. Random-bot logs from poke-env support pipeline debugging only. All scientific claims concern belief inference and generalization, with CPS as an empirical testbed.

3.4 Action Space

At each decision point the agent may use a move targeting an opponent slot or switch to a bench Pokémon. Actions are encoded as discrete classes in an action vocabulary built from training data. The exact action cardinality depends on format rules and logged move names.

3.5 Data Pipeline Status

The implementation roadmap proceeds in phases.

1. Local Showdown server and baseline logging (**complete**).
2. 13-token turn encoder and vocabulary (**complete**).

3. TurnTransformer behavioral cloning (**complete**).
 4. Metamon download and bridge to token format (**in progress**).
 5. Belief label extraction and multi-task training (**in progress**).
 6. Combinatorial split tooling and evaluation harness (**planned**).
- Appendix [A](#) lists repository paths for each component.

Chapter 4

Representation and Belief Targets

4.1 Design Goals

The observation representation must satisfy four constraints.

1. Fixed structure each turn so batching and attention masks are simple.
2. Explicit separation of field state, player party, and opponent party.
3. Compatibility with first-person partial observations from Metamon.
4. Discrete tokens suitable for embedding lookup and vocabulary growth.

4.2 Thirteen-Token Turn Layout

Each turn is encoded as exactly thirteen discrete tokens.

Table 4.1: Turn token layout. Indices are zero-based.

Index	Role
0	Field (weather, terrain, hazards on both sides)
1–6	Player party slots (padded to six)
7–12	Opponent party slots (partial info when unrevealed)

The field token concatenates sorted effect keys for weather, terrain, and side conditions. Each party token encodes species, discretized HP bucket, status, active flag, and fainted flag. Unfilled slots use a padding token.

A battle is a sequence of turns. Training samples use a causal context window of N prior turns plus the current turn. Left padding aligns the most recent turn to the final sequence position.

4.3 History Window

Given turn frames $x_1, \dots, x_t$, we form context

$$C_t = \text{pad}(x_{t-N+1}, \dots, x_t)$$

with a binary turn mask marking valid versus padded positions. The model may not attend to future turns. This implements a finite-history approximation to belief updating in a POMDP.

4.4 Belief Supervision Targets

Belief heads predict hidden opponent attributes knowable from full logs but not from the first-person view at time t . Primary targets in priority order are

1. species identity for each unrevealed opponent slot,
2. moves for revealed or partially revealed opponent Pokémon,
3. held item when inferable before reveal,
4. ability when not yet shown,
5. terastallization type in Gen 9 OU.

Each target is a multi-class prediction per opponent slot. Revealed or empty slots are masked out of the loss. Labels are extracted from Metamon full-state annotations at each turn boundary.

For species belief we define a vocabulary over base species names. Slot $s \in \{0, \dots, 5\}$ maps to opponent token index $7 + s$ and receives a label $y_{t,s} \in \mathcal{V}_{\text{species}} \cup \{\emptyset\}$ where $\emptyset$ denotes masked slots.

4.5 Combinatorial Keys for Evaluation

We define two keys for distribution shift experiments.

Team key Sorted tuple of six species names for the opponent roster.

Set key Tuple of species, item, ability, and move quartet for a slot.

Held-out keys define test battles where the opponent configuration never appeared in training. This tests compositional generalization beyond random battle splits.

4.6 Implementation

The `TurnTokenizer` class in `src/pokemon_thesis/tokenizer/` implements encoding from JSON battle records. Encoded trajectories are stored as JSONL with `token_ids`, `action`, and optional `belief` fields for multi-task training.

Chapter 5

Model and Training

5.1 Overview

We extend the existing `TurnTransformer` backbone with auxiliary belief heads and optional belief fusion into the policy head. The design keeps parameter count modest (on the order of one to five million) for Master’s-scale compute.

5.2 Hierarchical Encoder

5.2.1 Within-Turn Self-Attention

For each turn t , token embeddings are summed with slot position embeddings and passed through a transformer encoder over the thirteen slots:

$$h_t = \text{MeanPool}(\text{SlotEnc}(E(x_t) + P_{\text{slot}})).$$

Mean pooling over slot outputs yields one turn embedding $h_t \in \mathbb{R}^d$.

5.2.2 Cross-Turn Causal Attention

Turn embeddings receive turn position embeddings and pass through a causal transformer encoder:

$$\tilde{h}_t = \text{TurnEnc}_{\text{causal}}(h_{1:t} + P_{\text{turn}})_t.$$

The representation at the last valid turn index is used for prediction heads.

This two-level structure lets slot attention relate field state to individual party members within a turn while turn attention integrates evidence across time without peeking at future observations.

5.3 Policy Head

The policy head is a linear classifier over the final turn embedding:

$$\pi(a \mid C_t) = \text{softmax}(W_{\pi} \tilde{h}_{t^*}),$$

where t^* is the index of the most recent valid turn. Training minimizes cross-entropy against the human action (behavioral cloning).

5.4 Belief Heads

For each belief target type k and opponent slot s , a small MLP maps $\tilde{h}_{t*}$ to logits over the target vocabulary:

$$\hat{z}_{t,s}^{(k)} = \text{softmax}(f_s^{(k)}(\tilde{h}_{t*})).$$

Species prediction uses six slot-specific heads sharing the same trunk representation. Masked slots exclude terms from the belief loss.

5.5 Belief Fusion Variant

In an ablation we concatenate belief head softmax outputs to the policy input:

$$\pi(a \mid C_t) = \text{softmax}(W_\pi[\tilde{h}_{t*}; \hat{z}_t]).$$

This tests whether explicit belief features help action selection beyond shared representation learning in the trunk.

5.6 Training Procedure

5.6.1 Stage A: Belief-Augmented Behavioral Cloning

The multi-task loss is

$$\mathcal{L} = \mathcal{L}_{\text{action}} + \lambda_{\text{bel}} \sum_k \mathcal{L}_{\text{bel}}^{(k)}, \quad (5.1)$$

where $\mathcal{L}_{\text{action}}$ is action cross-entropy and each $\mathcal{L}_{\text{bel}}^{(k)}$ is masked cross-entropy over opponent slots. The weight λ_{bel} is tuned on validation data.

5.6.2 Stage B: Conservative Offline RL Fine-Tuning

We initialize from Stage A and optionally fine-tune with CQL or AMAGO on the same offline trajectories. No search or simulator rollouts run at inference.

5.6.3 Stage C: Evaluation

All benchmark battles use frozen checkpoints. No online weight updates occur during evaluation.

5.7 Baselines

Belief baselines include gradient-boosted trees and logistic regression on hand-crafted features, MLP on the latest turn only, GRU over turn embeddings, TurnTransformer without belief loss, and a flat causal transformer over $13 \times T$ tokens.

Policy baselines include BC with MLP, BC with GRU, BC TurnTransformer, belief-augmented TurnTransformer, and belief-augmented models with offline RL fine-tuning.

5.8 Architecture Summary

The `BeliefTransformer` class implements belief heads and optional fusion. Training is invoked via `scripts/train_belief_bc.py`.

Table 5.1: Default hyperparameters for the belief transformer.

Hyperparameter	Default
Model dimension d	128
Attention heads	4
Slot encoder layers	1
Turn encoder layers	2
Context turns N	8
Dropout	0.1
Optimizer	AdamW, lr 3×10^{-4}
λ_{bel}	tuned on validation

Chapter 6

Experimental Setup

6.1 Dataset

The primary corpus is Metamon Gen 9 OU human replays. The scale target is at least fifty thousand battles after filtering corrupted or extremely short games. Observations use first-person reconstructed from spectator logs. Actions come from human moves. Belief labels come from full-state opponent annotations.

Random-bot logs from poke-env support pipeline validation before human data is available at full scale.

6.2 Data Splits

We report three complementary split strategies.

6.2.1 Split A: Random Battle Split

Standard 80/10/10 train, validation, and test partition by battle ID. This measures in-distribution imitation and belief accuracy.

6.2.2 Split B: Held-Out Team Cores

Each opponent team receives a canonical key from sorted species names. Teams in a designated held-out set never appear in training. Test battles use opponents whose team key was unseen during training.

6.2.3 Split C: Held-Out Set Combinations

We hold out specific tuples of species, item, ability, and move quartet. This tests fine-grained combinatorial generalization beyond species identity alone.

6.2.4 Split D: Temporal Split (Optional)

When metadata permits, train on an earlier ladder period and test on a later period to measure temporal drift.

6.3 Belief Metrics

For each hidden-attribute type and model we report

- top-1 accuracy on masked unrevealed slots,
- top-3 accuracy where label space is large,

- expected calibration error (ECE),
 - Brier score for multi-class predictions,
 - reveal-turn anticipation (turns before official reveal when top-1 becomes correct).
- We plot belief accuracy versus turn number to show uncertainty narrowing over time.

6.4 Policy Metrics

Policy evaluation uses

- BC cross-entropy and top-1 action match versus human actions,
- offline policy value estimates when Q-learning is used,
- live win rate on poke-env against fixed opponents,
- generalization gap (in-distribution metric minus out-of-distribution metric).

The generalization gap is the primary summary statistic for RQ3.

6.5 Ablations

Table 6.1: Planned ablation studies.

Ablation	What it tests
Context $N \in \{4, 8, 16, 32\}$	History length for belief
Remove within-turn attention	Value of 13-token structure
Bidirectional turn encoder	Future leakage effects
$\lambda_{\text{bel}} = 0$ vs tuned	Necessity of belief loss
Belief fusion vs shared trunk	Explicit belief features
CQL vs AMAGO vs BC-only	Offline RL contribution

6.6 Statistical Reporting

We report mean $\pm$ 95% bootstrap confidence intervals over battles. Paired comparisons use the same test battles for each model. Win-rate significance uses paired bootstrap or McNemar tests where appropriate.

6.7 Minimum Viable Result

The thesis meets its pass bar if at least two of the following hold on combinatorial splits.

1. Belief-augmented TurnTransformer beats classical and recurrent belief baselines on aggregate hidden-species prediction.
2. Belief-augmented BC beats BC-only TurnTransformer on OOD team metrics with a smaller generalization gap.
3. Offline RL fine-tuning improves live win rate over BC-only without search at inference.
4. Ablations confirm causal history modeling contributes beyond flat sequence baselines.

6.8 Compute and Reproducibility

Training runs on a single GPU when available with CPU fallback. Checkpoints, vocabularies, and split manifests are saved under `checkpoints/` and `logs/`. The thesis PDF is built from `thesis/main.tex`.

Chapter 7

Results

7.1 Overview

This chapter reports behavioral cloning and combinatorial generalization results on Metamon Gen 9 OU human replays. Phase 1 established a turn-only **TurnTransformer** baseline on a large encoded subset. Phase 2 exported a belief-enriched corpus, trained three ablations on the same vocabulary, and evaluated them under random battle holdout (Split A) and held-out opponent team keys (Split B). All Phase 2 models use an eight-turn context window, $d = 128$, and twenty training epochs unless noted otherwise.

Corpus statistics for the scaled belief export are 69 411 turns from 2 000 battles, with 59 254 supervised move slots and thousands of item, ability, and terastallization labels on the active opponent. Gen 9 OU team preview reveals all opponent species at battle start, so species belief supervision is empty by design. Active-opponent move, item, ability, and tera type are the primary belief targets.

7.2 Phase 1 Baseline (Full Subset)

The turn-only **TurnTransformer** was trained on `logs/encoded_metamon_subset.jsonl` (approximately 255 000 turns). On a random 20% validation split this model reached **18.4%** top-1 action match rate across 988 action classes (checkpoint `checkpoints/human_bc/bc_best.pt`). This exceeds the chance rate and confirms that the tokenization pipeline and hierarchical encoder learn non-trivial structure from human trajectories alone.

7.3 Policy Quality (RQ2)

7.3.1 Random Battle Holdout (Split A)

Table 7.1 reports validation action accuracy on the belief corpus after training on 80% of battles. All three Phase 2 models share the same encoded file and token vocabulary.

Table 7.1: Top-1 action match on random validation battles (Split A), belief corpus, 2 000 battles, 20 epochs.

Model	Action classes	Val. accuracy
TurnTransformer BC (<code>human_bc_belief</code>)	425	20.4%
Matchup BC (<code>matchup_bc</code>)	425	20.5%
ActiveBelief (<code>active_belief</code>)	425	20.2%
ActiveBelief + fusion	425	19.6%

Turn history alone remains the strongest signal for action prediction at this scale. Explicit team matchup conditioning and auxiliary belief losses did not improve random-holdout action accuracy relative to the turn-only baseline. The best checkpoint for `active_belief` reached 20.2% validation action accuracy while simultaneously training move, item, ability, and tera heads.

7.3.2 Live Evaluation

Live win-rate experiments on `poke-env` against `RandomPlayer` and frozen BC opponents are planned but not yet included in this draft. **[TODO: Insert live battle win rates from `scripts/run_bc_match.py`.]**

7.4 Belief Inference (RQ1)

7.4.1 Active-Opponent Move Belief

The `ActiveBeliefTransformer` predicts four move slots plus item, ability, and tera type for the currently active opponent. Table 7.2 summarizes move-slot top-1 accuracy on the validation split.

Table 7.2: Move-slot belief accuracy on the validation split (2 000-battle corpus).

Model	History	Val. move belief
Linear probe (frozen BC trunk)	8 turns	10.5%
GRU baseline	8 turns	17.2%
ActiveBelief transformer	8 turns	28.1%
MLP baseline	latest turn only	35.5%

Scaling the export from 500 to 2 000 battles raised `active_belief` move belief from 15.5% to 28.1%. Classical baselines on the same labels show that a latest-turn MLP *outperforms* the transformer (35.5% versus 28.1%), while a GRU over eight turns reaches only 17.2%. Linear probes on the frozen BC trunk remain at 10.5% (Table 8.1). Much move information is therefore readable directly from the current turn tokens; deep causal history provides diminishing returns for this target at current scale.

7.4.2 Species and Set-Combination Beliefs

Species belief labels are unavailable in Gen 9 OU because team preview exposes all six opponent species before turn one. Split C (held-out item, ability, and move-set tuples) and calibration metrics (ECE, Brier score) are deferred to a later revision. **[TODO: Report item, ability, and tera belief accuracy and Split C results.]**

7.5 Combinatorial Generalization (RQ3)

7.5.1 Held-Out Opponent Teams (Split B)

Each opponent team is keyed by sorted species names. Training battles cover 1 615 team keys. Evaluation compares turns from battles whose opponent team appeared in training (*ID train teams*) against turns from 385 held-out team keys never seen during training (*OOD teams*). The generalization gap is ID train team accuracy minus OOD team accuracy.

Table 7.3 shows results at pilot (500 battles) and scaled (2 000 battles) corpus sizes.

Table 7.3: Action match under held-out opponent team keys (Split B).

Model	Scale	ID train	OOD teams	Gap
TurnTransformer BC	500	43.3%	41.4%	1.9%
Matchup BC	500	31.1%	29.3%	1.8%
TurnTransformer BC	2 000	45.2%	45.1%	0.05%
Matchup BC	2 000	42.6%	42.1%	0.4%
ActiveBelief	2 000	44.3%	44.4%	-0.03%
ActiveBelief + fusion	2 000	43.8%	43.3%	0.5%

On the pilot corpus, all models showed a one to two point gap between ID and OOD team splits. After scaling to 2000 battles, the gap collapsed to near zero for every architecture. Turn-only BC retained the highest absolute OOD accuracy (45.1%). Matchup conditioning did not close the gap further and scored lower in absolute terms. ActiveBelief matched turn-only OOD performance despite multi-task belief training.

Random Split A accuracy ($\approx 20\%$) is much lower than Split B team-conditioned accuracy ($\approx 42\text{--}45\%$) because the splits measure different quantities. Split A holds out entire battle trajectories. Split B evaluates action match on turns grouped by whether the opponent roster was seen during training, which partially aligns train and test action distributions when species composition is the primary shift variable.

7.6 Architecture Ablations (RQ4)

Phase 2 implements three incremental ablations on a shared trunk.

- **TurnTransformer BC** uses only turn-history tokens.
- **Matchup BC** adds a `TeamMatchupEncoder` over both teams’ six species and fuses the result into the final turn embedding before the action head.
- **ActiveBelief** extends Matchup BC with move, item, ability, and tera prediction heads trained with masked cross-entropy.

At 2000 battles, matchup and belief modules did not outperform turn-only BC on either Split A or Split B. This suggests that, at current data and model scale, the thirteen-token turn representation already encodes most team and belief information useful for action matching. Explicit structure helps move inference (RQ1) more than policy metrics (RQ2).

7.6.1 Belief Fusion

We tested concatenating belief softmax outputs (four move slots, item, ability, tera) to the policy head input (Chapter 5). After twenty epochs, `active_belief_fusion` reached **19.6%** Split A validation action match and **27.6%** move belief, versus 20.2% and 28.1% for parallel heads without fusion. OOD team accuracy was 43.3% (fusion) versus 44.4% (without fusion). Feeding explicit belief probabilities into the action head did not improve imitation at this scale.

Flat 13T token transformers and context-window sweeps remain future work. **[TODO: Report flat-sequence and context-window ablations.]**

7.7 Stage B Offline RL

Conservative offline RL fine-tuning (CQL or AMAGO) initialized from belief- augmented BC has not been run. **[TODO: Insert Stage B offline RL results if pursued.]**

7.8 Summary

Table 7.4 consolidates the scaled Phase 2 numbers. Belief accuracy refers to move-slot top-1 prediction for `active_belief` only. OOD action match is Split B OOD team accuracy.

Table 7.4: Results summary on the 2000-battle belief corpus (20 epochs).

Model	Move belief (val)	OOD action match	Gen. gap
TurnTransformer BC	—	45.1%	0.05%
Matchup BC	—	42.1%	0.4%
ActiveBelief	28.1%	44.4%	−0.03%

7.8.1 Relation to Hypotheses

H1 (belief inference). Mixed. The transformer learns non-trivial move beliefs and beats the GRU baseline (28.1% versus 17.2%), but a latest-turn MLP reaches 35.5% on the same split. History-conditioned attention is not required for strong move prediction when current-turn tokens already expose most clues.

H2 (belief usefulness). Not supported at current scale. Auxiliary belief loss did not improve action match over turn-only BC on either split.

H3 (combinatorial generalization). Mixed. Scaling data removed the OOD team gap for all models. Belief and matchup modules did not produce a smaller gap than the baseline.

H4 (offline RL). Not yet tested.

Reproducibility artifacts are stored under `checkpoints/` and `logs/ood_comparison.json`. The one-command replication script is `scripts/scale_belief_research.py`.

Chapter 8

Analysis and Interpretability

8.1 Overview

This chapter interprets the Phase 1 and Phase 2 results reported in Chapter 7. The analysis focuses on three patterns that emerged from scaled experiments on Metamon Gen 9 OU human replays

1. move-belief accuracy improves with data scale while policy metrics plateau across architectures,
2. opponent-team generalization gaps shrink with corpus size for every model, and
3. turn-only behavioral cloning matches or exceeds matchup- and belief-augmented variants on action prediction.

8.2 Scaling Belief Versus Scaling Policy

The `ActiveBeliefTransformer` increased validation move-slot accuracy from 15.5% on a 500-battle pilot export to 28.1% on a 2000-battle export, while validation action accuracy rose from 8.5% to 20.2% over the same period. The belief head therefore benefits directly from additional labeled trajectories.

Action accuracy for all three Phase 2 models converged near 20% on random battle holdout (Split A) after scaling, with turn-only BC slightly ahead. Auxiliary belief losses add parameters and optimization pressure but did not produce a policy advantage at $d = 128$ and twenty epochs. Two explanations are consistent with the data.

Shared trunk sufficiency. The hierarchical turn encoder may already compress history into an embedding useful for both action and hidden-set prediction. Explicit belief heads learn to *decode* that embedding for supervised targets without changing the trunk enough to improve actions.

Task interference. Multi-task training divides gradient signal among move, item, ability, and terra heads. If λ_{bel} is not carefully tuned, belief terms may slow action learning without improving generalization. Future work should sweep λ_{bel} and try belief fusion into the policy input (Chapter 5) rather than parallel heads only.

8.3 Why the Team Generalization Gap Collapsed

On the 500-battle corpus, Split B showed a 1.8–1.9 point gap between ID train teams and OOD teams. At 2000 battles the gap fell below half a point for every architecture, including turn-only BC.

This pattern suggests that the pilot gap was driven largely by *data coverage* rather than architectural failure. With only 500 battles, many held-out team keys had few or no nearby

examples in embedding space. Expanding to 2000 battles increased overlap in species co-occurrence patterns and move distributions, so OOD team turns became easier to match even without explicit matchup modules.

Matchup BC did not outperform turn-only BC on Split B at scale. The `TeamMatchupEncoder` pools six species embeddings per side and adds them to the final turn vector. Much of that information is already present in the opponent party tokens within each turn’s thirteen-token block. The explicit module may therefore be redundant given current tokenization.

Split B accuracy ($\approx 42\text{--}45\%$) exceeds Split A accuracy ($\approx 20\%$) because the metrics condition on different splits. Split A withholds entire battles. Split B groups turns by roster key while many action contexts remain familiar. Reporting both is essential to avoid over-interpreting high team-split numbers as solved imitation.

8.4 Linking Beliefs to Actions

The original plan correlated per-battle belief accuracy with per-battle action match on held-out teams. Species belief was unavailable in Gen 9 OU because team preview reveals all opponent species. Move-belief correlation analysis is still informative.

At scaled training, `active_belief` reached 28.1% move belief and 44.4% OOD team action match, while turn-only BC reached 45.1% OOD action match with no belief head. Higher move-belief accuracy therefore did not coincide with better OOD policy performance in this run.

Linear probes on frozen `human_bc_belief` trunk embeddings tell a sharper story (Table 8.1, script `scripts/probe_trunk_belief.py`). A fifteen-epoch linear probe on the final turn embedding reached only **10.5%** validation move-slot accuracy, versus **27.7%** for `active_belief` move heads on the same split. The BC trunk therefore does *not* expose hidden moves through an easy linear readout. Explicit belief heads and multi-task training learn representations that decode move labels far better than post-hoc probing.

Action accuracy nevertheless remains similar across models. Belief-specific features appear to live in a subspace that the shared action head does not exploit, or multi-task training reshapes the trunk in ways that help belief decoding without improving the action logit channel used at inference.

Table 8.1: Move-slot belief accuracy on validation battles (same split).

Method	Val. move belief
Linear probe on frozen BC trunk	10.5%
ActiveBelief move heads	27.7%

8.5 Implicit Belief in the Turn-Only Trunk

Turn-only BC and ActiveBelief achieved nearly identical OOD team accuracy (45.1% versus 44.4%) despite a large difference in supervised move-belief accuracy. Linear probes on the frozen BC trunk reached only 10.5% move belief (Table 8.1), so implicit linear encodings of unrevealed moves are weak. Explicit belief heads add substantial decodable structure without translating into better action match at current scale.

The thirteen-token layout exposes opponent party state, field conditions, and revealed moves each turn. Causal attention over eight turns integrates switch patterns, damage clues, and item triggers. That design aligns with the thesis motivation for structured partial observability without a separate belief filter.

8.6 Gen 9 OU and Belief Target Selection

Team preview fundamentally changed the belief task relative to older formats where species were hidden. The implementation pivot to active-opponent move, item, ability, and tera labels was necessary for non-empty supervision.

Species-centric hypotheses (H1 as originally framed) do not apply verbatim to this benchmark slice. Move-belief results should be read as evidence that history- conditioned transformers infer unrevealed moves, not that the full hidden-state posterior is solved.

8.7 Decision Tree and Attention Studies

Shallow decision trees on hand-crafted turn features and attention visualizations were planned to compare interpretable clues against neural errors. **[TODO: Train decision trees on damage hints, hazard state, and revealed moves; report feature importances versus transformer move-belief errors.]**

[TODO: Qualitative case studies on turns where move belief becomes correct before official reveal; inspect slot attention weights.]

Expected error modes remain relevant for future case studies

- rare move sets on held-out teams,
- ambiguous damage ranges before stat clues accumulate,
- deliberate unconventional sets underrepresented in training,
- late-game turns with truncated history relative to the eight-turn window.

8.8 Takeaways for RQ1–RQ4

RQ1. Transformers beat GRU on move belief but not a latest-turn MLP, which reaches 35.5% validation accuracy versus 28.1% for `active_belief`.

RQ2. Belief augmentation did not improve action match over turn-only BC at current scale.

RQ3. Combinatorial team shift was mitigated primarily by scaling data, not by matchup or belief modules.

RQ4. Hierarchical turn structure appears sufficient for strong BC at this scale. Belief fusion into the policy head did not improve action match over parallel heads alone (19.6% versus 20.2% Split A). Flat-sequence ablations remain open.

These interpretations motivate the limitations and future directions in Chapter 9 and the revised conclusions in Chapter 10.

Chapter 9

Limitations and Future Work

9.1 Domain Specificity

Results are measured on competitive Pokémon. Transfer to other partially observed simultaneous-action domains requires re-tokenization and new belief target definitions. The methods are domain-agnostic but empirical claims are benchmark-specific.

9.2 Label and Reconstruction Noise

Metamon first-person reconstruction may contain errors on edge-case mechanics. We mitigate with manual spot checks, corrupted-battle filtering, and comparison of random-bot versus human statistics.

9.3 Belief Labels After Reveal

Belief metrics focus on pre-reveal turns with masked slots. Including post-reveal slots would inflate accuracy trivially.

9.4 Compute and Model Scale

We intentionally use modest model size for Master’s-scale hardware. Scaling width, depth, and context length may narrow gaps to large black-box sequence policies [Grigsby et al., 2025] but is left for future work.

9.5 Offline RL Stability

CQL and AMAGO fine-tuning can be unstable on sparse or noisy data. Stage B is optional for thesis completion. Strong Stage A belief and BC results suffice for the minimum viable bar.

9.6 Gen 9 Team Preview

In Gen 9 OU, all opponent species are visible after team preview. Species-based belief and Split B keyed only by species composition may understate harder combinatorial shift involving unrevealed items, abilities, and move sets. Split C and active-opponent belief labels partially address this gap; full set- combination held-out evaluation remains open (Chapter 7).

9.7 Future Directions

- Self-confirming belief tokens inspired by [Zhang et al. \[2024\]](#) for opponent intent before move reveal.
- Variational belief modules with amortized inference over team identity.
- Cross-format evaluation (other generations or tiers) for broader generalization claims.
- Release of combinatorial split manifests as a standalone benchmark contribution.

Chapter 10

Conclusion

Partially observed simultaneous-action games require agents to infer hidden opponent state from interaction history and act without lookahead search. This thesis studied whether a causal transformer with structured thirteen-token turns, auxiliary belief supervision, and combinatorial held-out evaluation improves offline imitation learning in that setting. Competitive Pokémon Gen 9 OU human replays via Metamon served as the benchmark.

10.1 Contributions

The work delivers four concrete contributions.

Pipeline and representation. A reproducible path from Metamon replays to encoded turn sequences, belief-enriched labels, and `poke-env` evaluation hooks. Each turn is tokenized as thirteen discrete slots with causal history over $N = 8$ prior turns.

Model family. A shared `TurnTransformer` trunk with incremental ablations: turn-only BC, matchup-conditioned BC (`TeamMatchupEncoder`), and `ActiveBeliefTransformer` with move, item, ability, and terra heads.

Evaluation protocol. Complementary splits for random battle holdout (Split A) and held-out opponent team keys (Split B), with generalization gap as the primary combinatorial statistic. A one-command scale script reproduces export, training, and OOD evaluation.

Empirical findings on human data. Phase 1 BC reached 18.4% action match on $\approx 255\,000$ turns (988 action classes). Phase 2 on 2000 battles showed 20.4% random-holdout action match for turn-only BC, 28.1% move-belief accuracy for the belief-augmented model, and near-zero OOD team gap after scaling.

10.2 Summary of Findings

Belief inference and policy imitation responded differently to scale. Move-belief accuracy more than doubled when the export grew from 500 to 2000 battles. Action accuracy improved for all models but did not rank belief-augmented training above turn-only BC. Matchup conditioning over full team rosters did not help, likely because species composition is already encoded in opponent party tokens and because Gen 9 OU reveals all species at team preview.

The opponent-team generalization gap shrank from about two percentage points on the pilot corpus to under half a point at scale for every architecture. Data coverage therefore explains much of the early gap. Explicit belief and matchup modules did not outperform the trunk on Split B OOD accuracy.

Species belief, as originally envisioned for partially hidden rosters, does not apply to this format. The empirical belief results concern active-opponent move sets and related hidden attributes.

10.3 Hypotheses Revisited

H1 received partial support: the transformer learns history-conditioned move beliefs well above pilot accuracy, though classical baselines are not yet reported.

H2 was not supported at current scale: auxiliary belief loss did not improve action match over BC-only training.

H3 was mixed: scaling reduced OOD degradation for all models; belief and matchup modules did not achieve a smaller gap than the baseline.

H4 (offline RL fine-tuning) remains future work.

The minimum viable bar in Chapter 6 asked for at least two of four criteria. This draft meets the bar on (i) demonstrating non-trivial belief learning on hidden moves and (iv) evidence that structured turn history contributes through strong BC relative to chance, but not on belief improving OOD policy or offline RL gains.

10.4 Broader Impact

The scientific contribution is not competitive Pokémon play. It is an empirical study of transformer belief inference and combinatorial generalization under offline imitation. The negative result on belief-to-policy transfer is informative for domains where auxiliary prediction tasks are tempting but trunk capacity may already absorb the relevant signal. Practitioners should validate whether explicit belief heads help policy metrics on their splits before adding multi-task complexity.

10.5 Future Work

Priority directions follow from Chapters 8 and 9.

- Belief fusion into the policy head and sweeps over λ_{bel} .
- Linear probes on BC-only trunk embeddings to test implicit belief.
- Held-out set-combination splits (Split C) for item, ability, and move tuples beyond species keys.
- Live win-rate evaluation on `poke-env` and classical MLP/GRU baselines for RQ1.
- Conservative offline RL (CQL or AMAGO) initialized from scaled BC checkpoints.
- Larger context windows, model width, and corpus scale toward the full Metamon subset used in Phase 1.
- Release of combinatorial split manifests as a standalone benchmark artifact.

Artifacts are available under `checkpoints/`, `logs/`, and `scripts/scale_belief_research.py` for replication and extension.

Bibliography

- Lili Chen, Kevin Lu, Aravind Rajeswaran, Kimin Lee, Aditya Grover, Michael Laskin, Pieter Abbeel, Aravind Srinivas, and Igor Mordatch. Decision transformer: Reinforcement learning via sequence modeling. In *International Conference on Machine Learning*, 2021.
- Jake Grigsby, Yuqi Xie, Justin Sasek, Steven Zheng, and Yuke Zhu. Metamon: Scaling pokemon battle learning with human data. In *NeurIPS Datasets and Benchmarks Track*, 2023.
- Jake Grigsby, Yuqi Xie, Justin Sasek, Steven Zheng, and Yuke Zhu. Human-level competitive pokémon via scalable offline reinforcement learning with transformers. In *Reinforcement Learning Journal*, 2025.
- Leslie Pack Kaelbling, Michael L. Littman, and Anthony R. Cassandra. Planning and acting in partially observable stochastic domains. *Artificial Intelligence*, 101(1-2):99–134, 1998.
- Aviral Kumar, Aurick Zhou, George Tucker, and Sergey Levine. Conservative q-learning for offline reinforcement learning. In *Advances in Neural Information Processing Systems*, 2020.
- Brenden M. Lake, Tomer D. Ullman, Joshua B. Tenenbaum, and Samuel J. Gershman. Building machines that learn and think like people. *Behavioral and Brain Sciences*, 40, 2018.
- Yali Liu et al. Towards offline opponent modeling with in-context learning. In *International Conference on Learning Representations*, 2024.
- Richard S. Sutton and Andrew G. Barto. *Reinforcement Learning: An Introduction*. MIT Press, 2 edition, 2018.
- Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N. Gomez, Lukasz Kaiser, and Illia Polosukhin. Attention is all you need. In *Advances in Neural Information Processing Systems*, 2017.
- Yifan Zhang et al. Self-confirming transformer for belief-conditioned adaptation in offline multi-agent reinforcement learning. In *International Conference on Learning Representations*, 2024.

Appendix A

Implementation Reference

A.1 Repository Layout

Table A.1: Key paths in the thesis codebase.

Path	Purpose
src/pokemon_thesis/tokenizer/	13-token encoding
src/pokemon_thesis/model/	TurnTransformer, BeliefTransformer
src/pokemon_thesis/data/	Datasets and combinatorial splits
src/pokemon_thesis/metamon_bridge/	Metamon state conversion
scripts/train_bc.py	Behavioral cloning training
scripts/train_belief_bc.py	Multi-task belief + action training
thesis/main.tex	This document

A.2 Building the PDF

From the repository root on Windows PowerShell:

```
cd thesis
pdflatex main.tex
bibtex main
pdflatex main.tex
pdflatex main.tex
```

Or run `scripts/compile_thesis.ps1` which executes the same sequence.

A.3 Training Commands

Encode baseline logs:

```
python scripts/encode_logs.py logs/baseline_*.jsonl
```

Train BC TurnTransformer:

```
python scripts/train_bc.py logs/encoded_*.jsonl --epochs 30
```

Train belief-augmented model when encoded records include belief labels:

```
python scripts/train_belief_bc.py logs/encoded_*.jsonl --epochs 30
```

A.4 TurnTransformer Forward Pass

Given token IDs `token_ids` of shape $(B, T, 13)$ and turn mask `turn_mask` of shape (B, T) , the model applies slot self-attention per turn, causal turn attention across history, and linear action logits over the final valid turn embedding.

BeliefTransformer adds per-slot species heads and optional fusion into the policy input when enabled.